

One Collector to Rule Them All

Unified Observability for PostgreSQL Platforms

Dirk Aumueller

2026-04-22

Agenda

- Why? - The reason for this talk!
- A typical PostgreSQL Setup
- Component Logs & Metrics
- Component Configuration
- OpenTelemetry
- Demo
- Limits & Disadvantages
- Key Takeaways

Proventa

- A Venitus Company
- Database Transformation Services
- Customers in the private & public sector
- Founder of Postgres Usergroup Frankfurt (Main)
- Member of OSBA



- Dirk Aumüller, Associate Partner
- Generalist, not specialist!
- 10+ years Postgres experience

Why? - The reason for this talk!

Unified Observability: Monitoring Postgres Anywhere with OpenTelemetry

Thursday, October 23 at 09:25–10:15



Yogesh Jain
EnterpriseDB



PostgreSQL Conference Europe 2025

Room: **Omega 2**

Level: **Beginner**

PostgreSQL is widely used across cloud, Kubernetes, and bare metal environments. But how can we monitor it efficiently across all these setups without vendor lock-in or high costs?


In this talk, we'll explore how OpenTelemetry helps achieve consistent and scalable observability for Postgres - no matter where it runs. We'll walk through how to collect metrics and logs, integrate them with open-source observability tools, and optimize monitoring for both performance and cost.

Expect live demos and practical insights on making PostgreSQL monitoring simple, efficient, and future-ready with OpenTelemetry.

Key Takeaways: - How OpenTelemetry brings consistent Postgres observability across cloud, kubernetes, and bare metal - How to collect telemetry data (metrics and logs) for effective Postgres monitoring - How to use open-source tools (Prometheus, Loki, Grafana) to observe Postgres - Tips to reduce observability costs while keeping full visibility - Live demo of Postgres monitoring in cloud-native environments

What about all the other components of a Postgres service?

A typical PostgreSQL Setup



Component Log Formats

Component	Plain Text	CSV	JSON
PostgreSQL	yes	yes	yes
PGAudit ^{1 2}	yes	no	no
Patroni	yes	no	yes
pgBackRest ³	yes	no	no
PGBouncer	yes	no	no

¹ Extension [pgauditlogtofile](#) offers CSV & JSON with a separate PGAudit logfile

² Request for JSON implementation: [Github issue 293](#)

³ Request for JSON implementation: [Github issue 2619](#)

Log Format Examples

```
1 # Postgres
2 LOG: starting PostgreSQL 15.3 (.Ubuntu 15.3-0ubuntu0.22.04.1) on x86_64-pc-linux-gnu, compiled by gcc (.Ubuntu
3 LOG: listening on IPv4 address "127.0.0.1", port 5432
4
5 # Patroni
6 {"asctime": "2026-04-16 12:27:40,863", "levelname": "INFO", "message": "no action. I am (node1), the leader v
7
8 # pgBackRest
9 2026-04-16 12:26:20.361 P00 INFO: new backup label = 20260416-122618F
10 2026-04-16 12:26:20.437 P00 INFO: full backup size = 22.9MB, file total = 978
11 2026-04-16 12:26:20.437 P00 INFO: backup command end: completed successfully (.1769ms)
12
13 # PGBouncer
14 2026-04-16 12:26:47.317 UTC [147] LOG C-0x564359c703e0: postgres/testuser@[::1]:50722 login attempt: db=postgres
15 2026-04-16 12:26:47.324 UTC [147] LOG S-0x564359c86438: postgres/testuser@127.0.0.1:5432 new connection to server
16
17 # etcd
18 {"level":"info","ts":"2026-04-16T12:26:12.384085Z","caller":"version/monitor.go:116","msg":"cluster version c
```

Component Metrics

- Postgres, pgBackRest, PGBouncer
 - No metrics endpoint
 - Available with 3rd party tools, e.g. `postgres_exporter`, `pgbackrest_exporter`, `pgbouncer_exporter`
- Patroni, etcd
 - Prometheus compatible HTTP-Endpoint

```
1 dca# curl http://patroni:8008/metrics
2 # HELP patroni_version Patroni semver without periods.
3 # TYPE patroni_version gauge
4 patroni_version{scope="pgcluster1",name="node1"} 040101
5 ...
6 dca# curl http://etcd:2381/metrics
7 # HELP etcd_cluster_version Which version is running. 1 for 'cluster_version' label with current cluster vers
8 # TYPE etcd_cluster_version gauge
9 etcd_cluster_version{cluster_version="3.6"} 1
10 ...
```


Component Configuration

```

1 # etcd.conf
2 ETCD_LOGGER="zap"
3 ETCD_LOG_OUTPUTS="/var/log/etcd/etcd.log"
4 ETCD_LOG_LEVEL="info"
5 ETCD_LOG_FORMAT="json"
6 ETCD_ENABLE_LOG_ROTATION="true"
7 ETCD_LOG_ROTATION_CONFIG_JSON='{ "maxsize": 100, "maxage": 1, "maxbackups": 7, "localtime": true, "compress":

```

```

1 # patroni.yml - Postgres Logging
2 bootstrap:
3   dcs:
4     postgresql:
5       parameters:
6         log_destination: 'jsonlog'
7         logging_collector: 'on'
8         log_directory: '/var/log/postgres'
9         log_filename: 'postgresql-%a.log'
10        log_file_mode: '0640'
11        log_truncate_on_rotation: 'on'
12        log_rotation_age: '1d'
13        log_rotation_size: '1GB'
14        log_min_messages: 'warning'
15        log_min_error_statement: 'error'
16        log_connections: 'on'
17        log_disconnections: 'on'
18        log_error_verbosity: 'verbose'
19        log_hostname: 'off'
20        log_line_prefix: '%m [%r] [%p]: [l-%l] %u@%d,app=%a,e=s'
21        log_statement: 'none'
22        log_timezone: 'UTC'

```

```

1 # patroni.yml - Patroni Logging
2 log:
3   type: json
4   dir: /var/log/patroni
5   mode: 0644
6   level: INFO

```

```

1 # pgbackrest.conf
2 log-path=/var/log/pgbackrest
3 log-level-console=info
4 log-level-file=info

```

```

1 # pgbouncer.ini
2 [pgbouncer]
3 logfile = /var/log/pgbouncer/pgbouncer.
4 log_connections = 1
5 log_disconnections = 1
6 log_pooler_errors = 1
7 log_stats = 0

```

OpenTelemetry


- **OpenTelemetry** (OTel) is an open source observability framework
- Developed by the Cloud Native Computing Foundation (CNCF)
 - Provides vendor-neutral SDKs, APIs, & tools
- Collects, processes, & exports telemetry data (logs, metrics, & traces) in a single, unified format
- Exports to any observability backend, e.g. like Jaeger & Prometheus

! Important

OpenTelemetry becomes the dominant observability telemetry standard in cloud-native applications and its adoption is considered critical.

OpenTelemetry Collector

- **OTel Collector** is a vendor-neutral implementation of how to receive, process & export telemetry data
 - Serves as unified agent / collector
- **OTel Collector Contrib** repository offers many receivers, processors & exporters



Demo Setup

- Rocky Linux 10 UBI container image
 - systemd (privileged)
 - etcd
 - PostgreSQL v18
 - PGAudit
 - Patroni
 - pgBackRest
 - Pgbouncer
 - OTEL Collector Contrib
- Configuration copied into the container image

```
1 # host
2 dca# git clone https://github.com/dirkcaumueeller/otel-poc.git
3 dca# git checkout -b pgconf
4 dca# cd otel-poc
5 dca# chmod +x init.sh
6 dca# ./init.sh
7 dca# docker exec -it node1 bash
8
9 # container
10 root# systemctl start etcd
11 root# systemctl start patroni
12 root# systemctl start pgbouncer
13 root# systemctl start otelcol-contrib
```



Caution

- Not production ready!
- No mounts / volumes used - data is ephemeral!
- Credentials: username = password

OTel Collector Binary

- Different installation methods available
 - Docker, rpm, binary, ...
- Configuration via YAML file(s)

```
1 dca# otelcol-contrib --version
2 otelcol-contrib version 0.150.1
3
4 dca# otelcol-contrib --help
5 Usage:
6   otelcol-contrib [flags]
7   otelcol-contrib [command]
8
9 Available Commands:
10  completion  Generate the autocompletion script for the specified shell
11  ...
12
13 dca# otelcol-contrib --config=/etc/otelcol-contrib/otelcol-config.yaml
```

OTel Collector Configuration

```
1 receivers:                                # Collect data from sources
2   postgresql:                             # Dedicated receiver for PostgreSQL metrics
3     endpoint: 127.0.0.1:5432
4     transport: tcp
5     username: otel
6     password: otel
7     collection_interval: 10s
8     databases:
9       - postgres
10
11 processors:                             # Modify or transform data
12   batch:
13     timeout: 10s
14
15 exporters:                             # Pull / push based export of data
16   prometheus:
17     endpoint: '0.0.0.0:8900'
18
19 service:                                # Configure & enable components based on the configuration found
20   pipelines:                             # Consists of a set of receivers, processors & exporters
21     metrics:                             # Pipeline to publish Postgres metrics at Prometheus compatible http-endpoint
22       receivers: [postgresql]
23       processors: [batch]
24       exporters: [prometheus]
```

OTel Collector Receivers

```

1 receivers:
2   file_log/pgbouncer:                                # Use file receiver
3     include: [/var/log/pgbouncer/pgbouncer.log]
4     start_at: end
5     operators:                                       # Perform "little" tasks
6       - type: regex_parser
7         regex: '^(?<timestamp>\d{4}-\d{2}-\d{2} \d{2}:\d{2}:\d{2}\.\d{3} [A-Z]{3}) \[(?<pid>\d+)\] (?<error_s
8         timestamp:
9           parse_from: attributes.timestamp
10          layout: '%Y-%m-%d %H:%M:%S.%f %Z'
11          severity:
12            parse_from: attributes.error_severity
13            mapping:
14              debug:
15                - DEBUG
16                - NOISE
17              info: LOG
18              warn: WARNING
19              error: ERROR
20              fatal: FATAL
21          - type: move                                # Move attributes in common JSON structure
22            if: attributes["pid"] != nil
23            from: attributes["pid"]
24            to: attributes["process.pid"]
25          - type: move
26            if: body != nil
27            from: body
28            to: attributes["log.record.original"]
29          # Must be last operator-action
30          - type: move

```

OTel Attributes

- OTel defines a standardised set of common terms which are called attributes
- **General attributes**, e.g.
 - server.address
 - server.port
- Specific attributes, e.g.
 - db.system.name
 - deployment.status
- **Attribute Registry**

OTel Processors

```
1 processors:
2   memory_limiter:           # Prevents out of memory situations on the collector
3     check_interval: 1s
4     limit_percentage: 10
5     spike_limit_percentage: 5
6   resourcedetection/system: # Detect resource information from the host
7     detectors: ["system"]
8     system:
9       hostname_sources: ["dns"] # Get the fully qualified domain name
10  resource:                  # Add static information about data's origin
11    attributes:
12      - key: service.name
13        value: pg-db
14        action: upsert
15      - key: db.system.name
16        value: postgresql
17        action: upsert
18      - key: deployment.environment.name
19        value: development
20        action: upsert
21  batch:                     # Creates batches to better compress data & reduction of outgoing connections
22    timeout: 10s
```

OTel Collector Services

```
1 service:
2   pipelines:
3     metrics:
4       receivers: [sqlquery/custom, postgresql]
5       processors: [memory_limiter, resourcedetection/system, resource, batch]
6       exporters: [prometheus]
7     metrics/2:
8       receivers: [sqlquery/pgbouncer]
9       processors: [memory_limiter, resourcedetection/system, resource, batch]
10      exporters: [prometheus/2]
11     metrics/3:
12       receivers: [sqlquery/pgbackrest]
13       processors: [memory_limiter, resourcedetection/system, resource, batch]
14       exporters: [prometheus/3]
15
16     logs:
17       receivers: [file_log/postgresql, file_log/patroni, file_log/pgbouncer, file_log/pgbackrest]
18       processors: [memory_limiter, resourcedetection/system, resource, batch]
19       exporters: [file/rotation_with_custom_settings]
```

Live Demo



<https://github.com/dirkcaumueller/otel-poc>


Disadvantages & Limits

- Resource overhead
 - Collector processes consume CPU & memory, especially at high throughput
- Configuration verbosity
 - YAML pipelines can become large & hard to manage
- Steep learning curve
 - Pipelines, processors, exporters, and receivers have many moving parts
- Debugging is hard
 - Troubleshooting dropped or malformed telemetry in the pipeline is non-trivial
- Version instability
 - Components (especially contrib receivers) vary in maturity

Key Takeaways

- OTEL provides vendor-neutral observability
 - Avoid lock-in to any single monitoring tool
- Unified telemetry (metrics, traces, logs) in one pipeline reduces Postgres tooling sprawl
- End-to-end visibility from app code through connection pool to database
- Faster mean time to detect and resolve incidents
- Scales with your infrastructure
 - Same approach works from single node to large Postgres HA clusters
 - Environment agnostic (cloud native, hybrid, traditional)

Thank you!



PROVENTA 